


$$p \leq \alpha$$

Testing Hypotheses





Are these fingernail lengths
extreme?



Testing Hypotheses

1. Hypotheses and Testing Philosophies
2. Introducing the Frequentist P-Value
3. Test Statistics and All That
4. You're Doing it Wrong (or are you)

What is a Hypothesis? (in a statistical sense)

- **Point:** Mean Fingernail Length == 5cm
- **Model:** If I increase Kelp, Biodiversity increases
- **Explanatory:** Fertilizers explain more variation in the data than Pesticides

Three Flavors of Statistical Inference

1. **Frequentist:** What's the chance we observed something like the data given our hypothesis?
2. **Likelihoodist:** What's the likelihood of our hypothesis relative to others?
3. **Bayesian:** How much do we believe our hypothesis given our data?

Frequentist Inference: the probability of observing this data, or more extreme data, given that a hypothesis is true

P-value
or
Confidence Interval


Our Hypothesis

‘Extreme’ fingernail lengths are not different from standard fingernail lengths

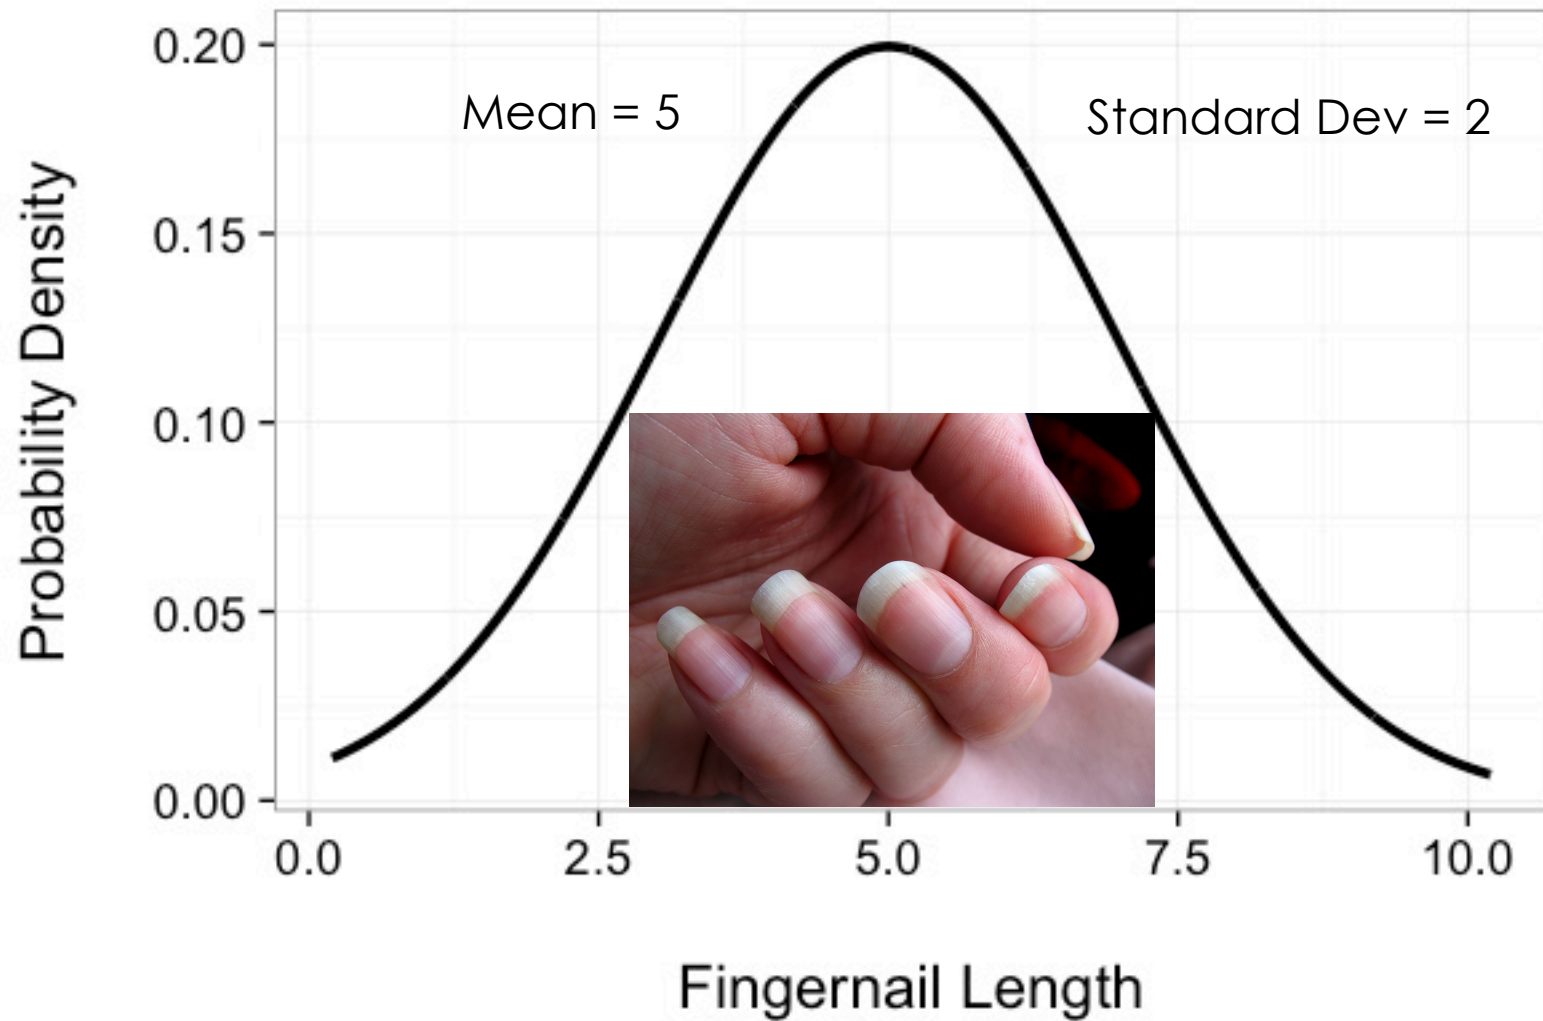
Testing Hypotheses

1. Hypotheses and Testing Philosophies
2. Introducing the Frequentist P-Value
3. Test Statistics and All That
4. You're Doing it Wrong (or are you)

Our Sample



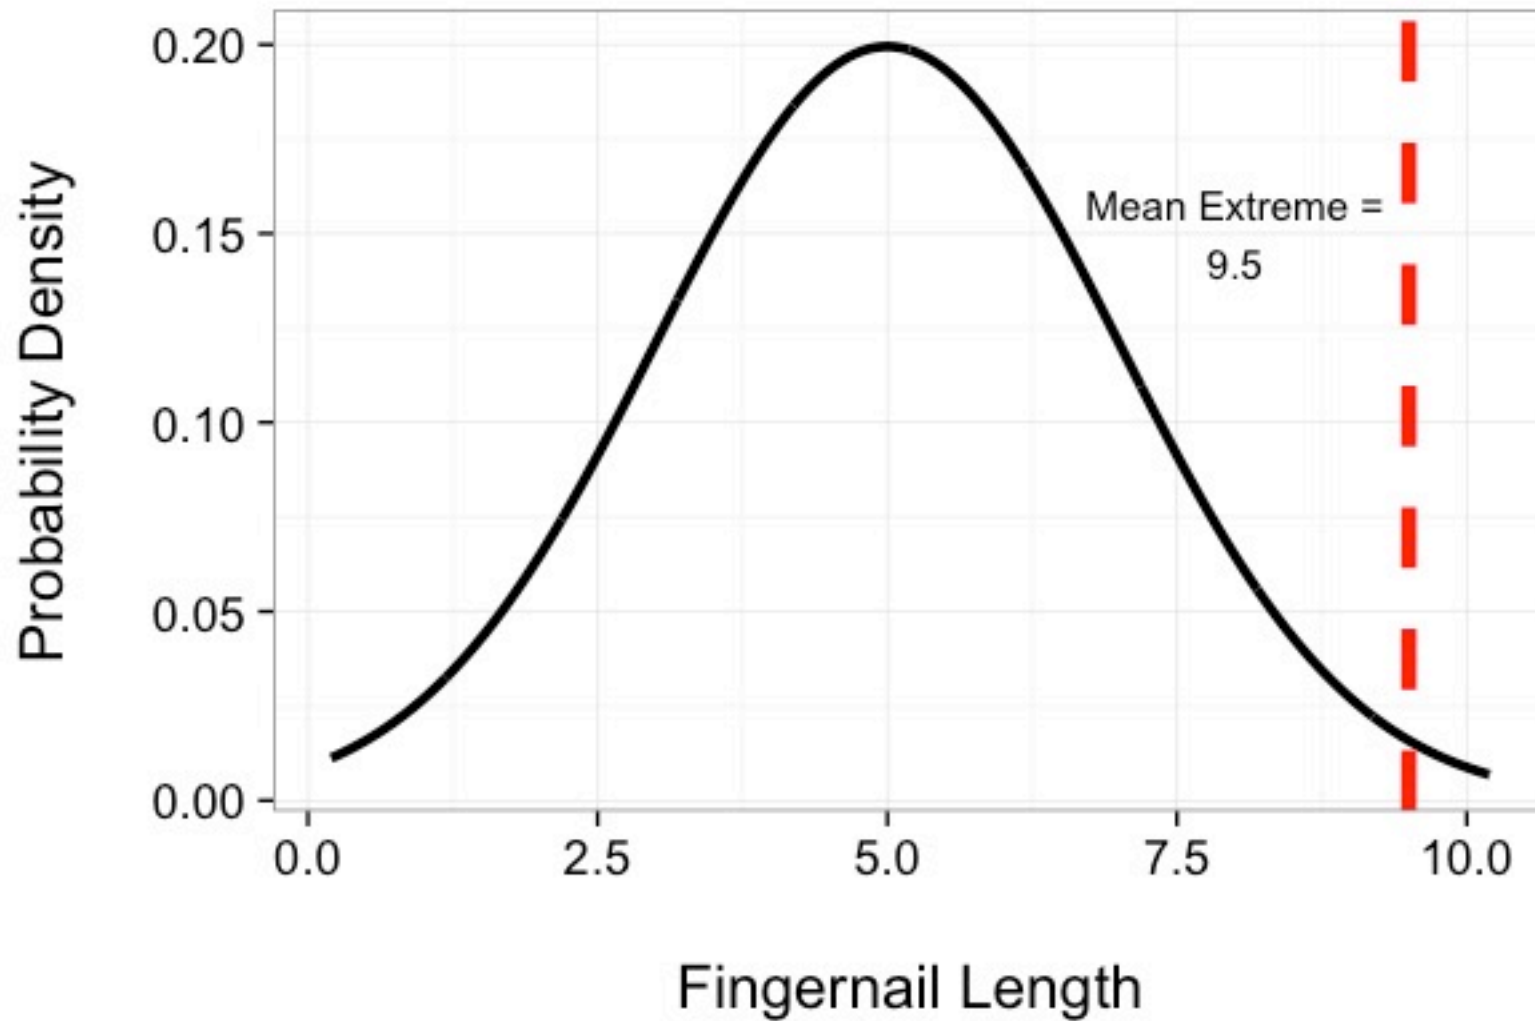
Assume a Distribution of Fingernail Sizes




Our Hypothesis



‘Extreme’ fingernail lengths are from a distribution with mean of 5 and SD of 2.

Is our size extreme?



What about even more extreme values?





Extreeeeeme!



P-Values: Fisher



R.A. Fisher

The probability of observing a value or more extreme value given a specified hypothesis.

Testing Hypotheses

1. Hypotheses and Testing Philosophies
2. Introducing the Frequentist P-Value
3. Test Statistics and All That
4. You're Doing it Wrong (or are you)

Null Hypothesis Testing

- Frequentist testing of whether ***something*** is different from a null expectation
- Example uses:
 - An estimate is not different from 0
 - The difference between two groups is not different from zero
 - A predictor provides no additional explanation of patterns in the data

Test Statistics: Making the World Sensible (and Null)


1. Create a null distribution
2. Use your data to calculate a test score
3. Calculate the p-value for your data in the context of that null distribution
 - $P(D | H_0)$:

Testing Fingernails Against a Normal

H_0 = Mean of sample is not different from the rest of the population

1. Assume a normal distribution with an SD of 2
2. Calculate the difference between the mean of our sample and 5 = 4.5
3. Assess the p-value of 4.5 against the normal distribution with mean 0

But...this is my data



The Arrival of a Test Statistic

H_0 = Mean of sample is not different from the rest of the population

1. Assume a normal distribution with an SD of 1 (standard normal curve)
2. Calculate the difference between the mean of our sample and 5 = 4.5
3. Divide by the Standard Deviation of the population and the square root of the sample size (assumed SE of a population Sample) – z score!

- Z-Score


$$z = \frac{\bar{x} - \mu}{\sigma / \sqrt{n}}$$

$$z = 4.5 / (2/\sqrt{30})$$


$$= 12.32$$



$$p(12.32 \mid \mu=0) < 0.0001$$

Z-Test



There Are Two Tails





Yes. These are super odd.





What do you do with a p-value?

- P-values give the evidence for the probability of your data, or more extreme data, given a hypothesis
- As the scientist, **you** decide whether it is grounds for rejecting a hypothesis
- In a frequentist framework, you can only reject a hypothesis – never `accept`

Testing Hypotheses

1. Hypotheses and Testing Philosophies
2. Introducing the Frequentist P-Value
3. Test Statistics and All That
4. You're Doing it Wrong (or are you)


How We Can Screw This Up

| Is the Hypothesis True or False? | | | |
|--------------------------------------|----------------------------|---|--|
| | | True | False |
| Test Result Against Hypothesis | Hypothesis Not Rejected |  | Type I error |
| | Hypothesis Rejected | Type II error |  |


Probability of a type I error = α

Probability of a type II error = β

Null Hypothesis Significance Testing



Egon Pearson



Jerzy Neyman

Problem: What is an acceptable α ?

Answer: 0.05. You have a 1 in 20 chance of committing a type I error

Problems with NHST

- A realistic alpha can depend on your study design
 - E.g., Large sample size = lower p value
- Ignores β
 - Tradeoff between α and β
- Conflation of scientific significance and statistical significance
 - File-drawer effect
- We are human
 - If $p \leq 0.05$ makes your career, you will do a lot to obtain it!